

RA_LAB3_ITDSIU21095

October 15, 2022

1 Problem 1

```
[1]: import pandas as pd, numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
import statsmodels.api as sm
import statsmodels.formula.api as smf
```

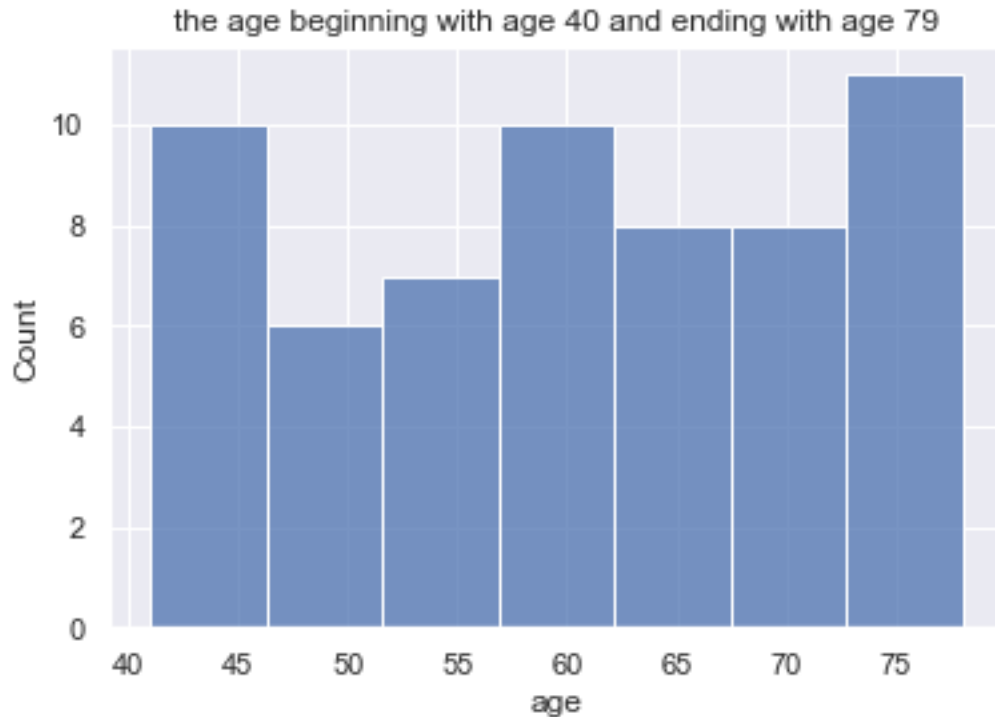
```
[15]: df = pd.read_csv('CH01PR27.txt', sep = '\s+', header = None, names = ['the_
    ↳measure of muscle mass', 'age'])
df.head()
```

```
[15]: the measure of muscle mass  age
0                                106  43
1                                106  41
2                                 97  47
3                                113  46
4                                 96  45
```

1.1 a.

Prepare a histogram for the ages X_i . What information does your plot provide? Is this plot consistent with the random selection of women from each 10-year age group? Explain.

```
[23]: sns.set_theme()
sns.histplot(data=df, x = 'age').set(title='the age beginning with age 40 and_
    ↳ending with age 79');
```



The information the plot provide is a nutritionist randomly selected 15 women from each 10-year age group, beginning with age 40 and ending with age 79. We can noticed that there are a slightly different distribution between these group. However this difference can be accepted and this plot can be consistent with the random selection of women from each 10-year age group.

1.2 b.

Obtain the residuals and prepare a normal probability plot of the residuals. Does the distribution of the residuals appear to be symmetrical?

```
[24]: X = df['age']
      Y = df['the measure of muscle mass']
      X_bar = np.mean(X)
      Y_bar = np.mean(Y)
      X_err = X - X_bar
      Y_err = Y - Y_bar
      print(X_bar, Y_bar)
      X_err.head()
```

```
59.983333333333334 84.96666666666667
```

```
[24]: 0    -16.983333
      1    -18.983333
```

```
2   -12.983333
3   -13.983333
4   -14.983333
Name: age, dtype: float64
```

```
[25]: A = np.sum(X_err*Y_err)
      B = np.sum(X_err**2)
      print(A, '\n', B)
```


```
-9771.033333333333
8210.983333333335
```

```
[26]: b1 = A / B
      b0 = Y_bar - b1*X_bar
      print(b1, b0)
```

```
-1.1899955141385823 156.34656425641265
```

```
[27]: n= len(X)
      Y_hat = b0 + b1 * X
      resid = Y - Y_hat
```

```
[34]: stats.probplot(resid, dist="norm", plot = plt);
      plt.show();
```




The distribution of the residuals appear to be symmetrical. However, the distribution in the right top corner and left bottom corner is not symmetrical.


1.3 c.

Plot the residuals against \hat{Y}_i and also against X_i on separate graphs to ascertain whether any departures from regression model (2.1) are evident. Do the two plots provide the same information? State your conclusions.

```
[40]: sns.residplot(x=Y_hat,y= resid, data=df);
```



```
[41]: sns.residplot(x=X, y=resid, data=df);
```



The two plots provide the same information. Additionally, there is an outlier in the left top corner in plot 1 and in the right top corner in plot 2

1.4 d.

Assume that (3.10) is applicable and conduct the Breusch-Pagan test to determine whether or not the error variance varies with the level of . Use $\alpha = 0.01$. State the alternatives, decision rule, and conclusion. Is your conclusion consistent with your preliminary findings in part (c)?

```
[43]: model = smf.ols('Y ~ X', data=df)
      results = model.fit()
      results.summary()
```

```
[43]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  Y      R-squared:                0.750
Model:                            OLS      Adj. R-squared:           0.746
Method:                 Least Squares      F-statistic:                174.1
Date:                Sat, 15 Oct 2022      Prob (F-statistic):       4.12e-19
Time:                  10:21:59      Log-Likelihood:           -210.17
No. Observations:                  60      AIC:                     424.3
Df Residuals:                      58      BIC:                     428.5
```

```

Df Model:                                1
Covariance Type:                        nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept    156.3466      5.512     28.363     0.000     145.313     167.381
X             -1.1900      0.090    -13.193     0.000     -1.371     -1.009
=====
Omnibus:                1.213   Durbin-Watson:                2.416
Prob(Omnibus):           0.545   Jarque-Bera (JB):           1.208
Skew:                    0.319   Prob(JB):                     0.547
Kurtosis:                2.726   Cond. No.                     319.
=====

```

Notes:

```

[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
"""

```

The alternative conclusions are: $H_0 : 1 == 0$

$H_a : 1 \neq 0$

The decision rule is: If $X^2 \leq X^2_{BP}$, conclude H_0

If $X^2 > X^2_{BP}$, conclude H_a

```

[46]: alpha = 0.01
      resid_square = resid**2
      model_1 = smf.ols('resid_square ~ X', data=df)
      results_1 = model_1.fit()
      results_1.summary()

```

```

[46]: <class 'statsmodels.iolib.summary.Summary'>
      """

```

```

                                OLS Regression Results
=====
Dep. Variable:          resid_square   R-squared:                0.074
Model:                  OLS           Adj. R-squared:           0.058
Method:                 Least Squares   F-statistic:              4.615
Date:                  Sat, 15 Oct 2022   Prob (F-statistic):       0.0359
Time:                  10:29:57          Log-Likelihood:          -349.29
No. Observations:        60             AIC:                   702.6
Df Residuals:           58             BIC:                   706.8
Df Model:                1
Covariance Type:        nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----

```

```
-----
Intercept    -53.5326    56.015    -0.956    0.343    -165.659    58.593
X              1.9690     0.917     2.148     0.036     0.134     3.804
=====
Omnibus:                66.573    Durbin-Watson:                2.197
Prob(Omnibus):           0.000    Jarque-Bera (JB):            547.395
Skew:                    3.051    Prob(JB):                    1.36e-119
Kurtosis:               16.481    Cond. No.                    319.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

"""

```
[48]: SSR_star= sum((results_1.fittedvalues - np.mean(resid_square))**2)
SSR_star
```

```
[48]: 31833.428951644546
```

```
[51]: X2_BP = (SSR_star/2) / ((SSE/n)**2)
X2_BP
```

```
[51]: 3.817124858502205
```

```
[82]: X2 = stats.chi2.ppf(q=1-0.01,df=1)
p =1- stats.chi2.cdf(x=3.817,df=1)
print(X2,p)
```

```
6.6348966010212145 0.05073500104800588
```


Conclusion: to control the alpha risk at 0.01, we require $X^2(0.99;1) = 6.6$. Since $X^2_{BP} = 3.817 \leq 6.63$, we conclude H_0 , that the error variance is constant.

2 Problem 2:

```
[102]: df1 = pd.read_csv('CH03PR18.txt', sep = '\s+', header =None,
    names=['hours','lot size'])
df1.head()
x= df1['lot size']
y= df1['hours']
x_new = np.sqrt(x)
```

2.1 a.

```
[72]: sns.scatterplot(x='lot size', y='hours', data=df1, color = 'black');
```



- A linear relation does not appear to be adequate here.
- The regression relation in the scatter-plot appears to be curvilinear here.
- The variability across different X levels appears to be fairly constant, thus a transformation to X is more appropriate here.

2.2 b.

```
[104]: model = smf.ols('y ~ x_new', data=df1)
results = model.fit()
results.summary()
```

```
[104]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  y    R-squared:                0.770
Model:                            OLS    Adj. R-squared:           0.768
Method:                 Least Squares    F-statistic:                365.7
Date:                Sat, 15 Oct 2022    Prob (F-statistic):       1.29e-36
Time:                  11:56:43    Log-Likelihood:          -232.88
No. Observations:                111    AIC:                     469.8
Df Residuals:                    109    BIC:                     475.2
Df Model:                          1
Covariance Type:                nonrobust
```


	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.2547	0.639	1.964	0.052	-0.012	2.521
x_new	3.6235	0.189	19.124	0.000	3.248	3.999
Omnibus:		1.123	Durbin-Watson:			2.098
Prob(Omnibus):		0.570	Jarque-Bera (JB):			1.160
Skew:		-0.151	Prob(JB):			0.560
Kurtosis:		2.600	Cond. No.			12.3

Notes:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

"""

- The estimated linear regression function for the transformed data is slightly higher than it original.

2.3 c.

```
[105]: sns.scatterplot(x=x_new, y= y, data =df1);
```



Conclusion: The scatter-plot shows a reasonably linear relation, the estimated regression line appears to be a good fit to the transformed data.

2.4 d.

```
[109]: x_bar = np.mean(x)
y_bar = np.mean(y)
x_err = x - x_bar
y_err = y - y_bar
print(x_bar, y_bar)


a = np.sum(x_err*y_err)
b = np.sum(x_err**2)
print(a, '\n', b)

b1_2= a / b
b0_2 = y_bar - b1_2*x_bar
print(b1_2, b0_2)


y_hat = b0_2 + b1_2 * x
y_hat_new = b0_2 + b1_2 * x_new
resid1 = y - y_hat_new

sns.residplot(x=y_hat_new, y= resid1,data=df1,color='red');
```

```
11.36936936936937 12.926486486486487
2608.7040540540543
4891.855855855856
0.5332749228355274 6.863486913347429
```



```
[110]: stats.probplot(resid1, dist="norm", plot = plt);
plt.show();
```



Conclusion The normal probability plot shows that points fall reasonably close to a straight line, with very small tails deviating from the line, suggesting that the distribution of the error terms is approximately normal. These two plots above show a good fit between model and data points. Q - Q plot show the distribution of the residuals appear to be symmetrical.

2.5 e.

2.5.1 $Y = 6.863486913347429 + 0.5333 X$